

INDUSTRY

- 2019–23 Data Scientist (Advanced Algorithms) **THE SIGNAL GROUP**, Athens, GR
Athens is a global shipping hub, with Greek shipping companies control a fifth of the world's fleet. Thenamaris (\$5B) offshoot Signal is building an algorithmic, data-driven platform, supporting 100+ firms.
Technical Product Manager: Led team of 7 data scientists & engineers, in modeling and product delivery of production systems for Natural Language Processing (NLP) entity extraction & time series forecasting. Direct report to Chief Strategy Officer & SVP of Advanced Algorithms. Client-facing with ship owners, brokers, quantitative analysts, & hedge-funds.
Time Series Forecasting: Modeled latent port waiting times globally, designed back-end system design to provide daily point-in-time updates for quantitative desks at shipping firms & hedge funds. Led technical sales meetings with experts.
NLP: Responsible for several explainable information extraction algorithms for unstructured commercial communications (Real-time processing 100k+ messages/day of cargoes, tonnage, fixtures, vessel specifications). Automated ship-brokers' daily information workflows, producing structured data for several major services: platform analytics & interfaces, cargo-to-ship matching, vessel forecasting, & commodity flows.
- 2018–19 Software Engineer **PLUME DESIGN, INC.**, Palo Alto, CA
Machine Learning: Home-network security model via clustering device DNS data with partially known threat levels. ML operational assistance, data pipeline migration (EC2, Python, R, Spark)
API & Front-End: Implemented security features (Node.js, MongoDB), diagnostic charts (Angular, D3.js, Highcharts), client-side user authentication, via test driven development (Mocha, Chai), for multiple services: pods, app, operation center.
- 2016 Software Developer **THE INSTITUTE FOR GENOME SCIENCES & SOCIETY AT TEXAS A&M**
gQTL: A Web Application for QTL Analysis Using the Collaborative Cross Mouse Genetic Reference Population [2018]. Designed web interface (RShiny, CSS, JS, SQL) for bioinformaticians, genetic researchers. Interface handled users, data visualization, preprocessing data, job management for QTL analysis.
† Published and spotlighted in, [G3: Genes, Genomes, Genetics](#)
- 2015 Software Developer **BAYLOR SCOTT & WHITE** and **TEXAS A&M HEALTH SCIENCE CENTER**
Data processing and analysis. Tabled patients' gene, isoform expression counts pre/post treatments ([public data](#)). Paired *t*-Tests, correlation analysis (under project statistician) for 20k+ elements (on [Github](#)).
† Acknowledged in [American Journal of Pathology](#) publication.

EDUCATION

- 2024– Ph.D. Statistics **ATHENS UNIVERSITY OF ECONOMICS & BUSINESS**
Thesis: Analysis of biomedical data via functional distributions. Methodology involving Gaussian & Neural Diffusion Processes with applications in heterogeneous epidemic models & survival curve extrapolation from clinical trials.
Supervisors: Nikos Demiris (AUEB & Cambridge), Petros Dellaportas (AUEB & University College London).
- 2022–24 M.Sc. Statistics
Thesis: Precision Medicine: Comparison of Outcome Weighted Learning estimators of Optimal Individualised Treatment Rules (ITR) for Right-Censored Survival Data. Bio-statistics, Causal inference, semi-parametric efficiency, statistical machine learning.
Supervisors: Giorgos Bakoyannis (Indiana University), Nikos Demiris (AUEB & Cambridge).
- 2014–19 B.Sc. Computer Science **NORTHWESTERN UNIVERSITY**
Graduate-Level Research **Planetary Scale Computing Group**
Google AMP improves web-page load time and other QoE metrics, but used only by news sites. In aim to democratise internet access to low-bandwidth areas, we evaluated automatic-AMP converters' across web at scale, designed and simulated experiments for QoE and DOM information loss.